

## Final Report

### Implementation of a Neural Network for Multi-Class Image Classification

#### **1. Introduction**

This project focuses on building and training a neural network from scratch to classify grayscale images of alphanumeric characters into 62 classes (digits 0-9, uppercase letters A-Z, and lowercase letters a-z).

#### **2. Dataset Preparation**

##### **2.1 Dataset Structure**

The dataset is organized into 62 folders (Sample001 to Sample062):

- Sample001 to Sample010: Represent digits 0-9.
- Sample011 to Sample036: Represent uppercase letters A-Z.
- Sample037 to Sample062: Represent lowercase letters a-z.

##### **2.2 Preprocessing**

The dataset is processed using the `load_and_preprocess_images` function, which performs the following steps:

1. **Image Loading:** Images are loaded using the Python Imaging Library (PIL).
2. **Grayscale Conversion:** Images are converted to grayscale for simplicity.
3. **Resizing:** All images are resized to 80x60 pixels to standardize input dimensions.
4. **Normalization:** Pixel values are scaled to the range [0, 1].
5. **Flattening:** Each image is flattened into a vector of size 4800 (80x60).
6. **Label Encoding:** `get_class_from_folder` Maps folder names to class labels.
  - 0-9 for digits.
  - 10-35 for uppercase letters.
  - 36-61 for lowercase letters.

##### **2.3 Data Splitting**

The `split_data` function is used to split the dataset into training and testing sets. For each class:

- 50 random samples are selected for training.
- 5 random samples are selected for testing.

This ensures a consistent 50:5 split while maintaining randomness.

### **3. Neural Network Architecture**

#### **Input Layer**

- Size: 4800 neurons, corresponding to the flattened image vector.
- Input Representation: Pixel intensities scaled between  $[0, 1]$ .

#### **Hidden Layer**

- Size: 100 neurons.
- Activation Function: Sigmoid.

#### **Output Layer**

- Size: 62 neurons, one for each class.
- Output Representation: A matrix of  $1 \times 62$  probabilities indicating the likelihood of each class.
- Activation Function: Sigmoid.

### **4. Training Procedure**

#### **4.1 Initialization**

The `initialize_parameters` function initializes weights and biases:

- **Weights:** Xavier initialization to stabilize signal flow.
- **Biases:** Initialized to 1.

#### **4.2 Forward Propagation**

The `forward_propagation` function computes activations for the hidden and output layers:

- Hidden layer activations are computed using the Sigmoid function.
- Output layer activations are computed using the Sigmoid function.

#### **4.3 Loss Function**

The `compute_cost` function calculates cross-entropy loss:

- Output activations ( $A_2$ ) are clipped to prevent numerical instability ( $\log(0)$ ).

#### **4.4 Backward Propagation**

The `backward_propagation` function computes gradients for weights and biases:

- Gradients are calculated for both layers using the derivative of the Sigmoid function.

#### **4.5 Gradient Clipping**

The `clip_gradients` function ensures gradient values remain within a specified range to prevent exploding gradients.

#### 4.6 Hyperparameters:

- Learning Rate: 0.01
- Epochs: 2000.
- Batch size: 32

#### 4.7 Steps Per Epoch:

- Shuffle the training data.
- Divide the training data into mini-batches.
- Perform forward propagation for each mini-batch.
- Compute the loss.
- Perform backward propagation to compute gradients.
- Update weights and biases using gradient descent in the `update_parameters` method.

### 5. Challenges and Optimizations

I faced the problem of **Vanishing Gradients** during initial training which is common when using Sigmoid as the activation function. Due to this the loss remained high ( $> 5$ ) and the accuracy on the training and testing dataset was low ( $< 3\%$ )

Several techniques were used to improve the model:

1. **Xavier Initialization:** Scaled weights based on input size to stabilize signal flow and prevent gradient saturation in sigmoid neurons. This ensured consistent gradient propagation and better convergence.
2. **Mini-Batch Processing:** Trained using random subsets of the dataset, stabilizing updates and improving convergence by reducing noise while maintaining stochasticity.
3. **Gradient Clipping:** Restricted the magnitude of gradients to avoid exploding gradients and ensure smoother, controlled updates.

All these measures significantly increased the accuracy of the model


### 6. Results

- **Training Accuracy:** 67.42%.
- **Testing Accuracy:** 43.55%.

True: 22  
Pred: 58




True: 23  
Pred: 46




True: 32  
Pred: 32




True: 48  
Pred: 53



True: 8  
Pred: 8




True: 0  
Pred: 56



True: 39  
Pred: 39



True: 47  
Pred: 21



True: 32  
Pred: 32



True: 34  
Pred: 29

